

# LLMs and prompting

Dr. Manuel Pita & Zuil Pirola



UNIVERSIDADE  
LUSÓFONA

# What is a Prompt?

A **prompt** is a text instruction that a user gives to a **language model (LLM)** to make it perform a useful task. When the model receives the prompt, it generates a response **token by token**, using the prompt as context.

The prompt acts as a **guide**, helping the model generate outputs aligned with the user's goal.

The process of designing, testing, and refining effective prompts is known as **Prompt Engineering**.

# Prompt Engineering

Prompt engineering is a set of techniques used to design effective ways of interacting with **large language models (LLMs)** and external tools, especially for complex tasks like **reasoning** and **problem-solving**.

It goes beyond simply writing prompts and includes understanding model capabilities, improving safety, and enhancing models with domain knowledge and external resources.

# Reasoning

**Reasoning** is the process by which a Large Language Model (LLM) generates logically structured responses by recognizing patterns in its training data and predicting sequences of tokens that follow coherent logical relationships.

Rather than “thinking” like a human, LLMs perform reasoning by using statistical patterns to simulate step-by-step logic, drawing connections between concepts, facts, and rules to produce structured and goal-oriented outputs.

# LLMs settings

When working with LLMs through an **API**, you can adjust parameters to control the **style, length, accuracy, and creativity** of model responses.

These settings help improve **reliability, relevance**, and **cost efficiency**, but usually require experimentation to optimize.

# LLMs settings

**Temperature** controls how *random* or *deterministic* the model's next token selection is.

It works by adjusting the probability distribution of possible next tokens:

- **Low Temperature (0.0–0.3)**

Narrows the probability distribution → the model consistently selects the most likely next token.  
Result: **factual, stable, predictable** responses.

- **Medium Temperature (0.4–0.7)**

Allows some variation while maintaining coherence.  
Result: **balanced** responses—useful for most tasks.

- **High Temperature (0.8–2.0)**

Widens the distribution → increases randomness and variety.  
Result: **creative, exploratory, but less reliable** outputs.

# LLMs settings

**Top-P** controls how many possible next tokens the model is allowed to consider when generating text.

Instead of sampling from all tokens, the model selects only from the smallest set of tokens whose cumulative probability reaches P (e.g., 0.9 = 90% of the probability mass).

- **Low Top-P (e.g., 0.1–0.3)**

The model picks from only the most likely tokens → **precise, focused, predictable** responses.

- **High Top-P (e.g., 0.8–1.0)**

The model considers many more tokens, including less probable ones → **more diverse, creative**, but sometimes less consistent output.

# LLMs settings

Both Temperature and Top-P control **randomness and diversity** in generation, but they do so in *different ways*. When used together:

- Their effects can **overlap**, making the output less predictable.
- It becomes harder to understand which parameter is influencing the behavior.
- Tuning becomes more complex and results less stable.

# LLMs settings

**Max Tokens** defines the **maximum number of tokens** the model can generate in a response.

Helps control:

- Response length
- **Relevance** (avoids overly long outputs)
- **Cost** (fewer tokens = lower usage)

A token can be a full word, part of a word, or punctuation, depending on the language.

In case of GPT use tiktoken.

# LLMs settings

```
import tiktoken

def count_tokens(text: str, model: str = "gpt-4o-mini"):
    """
    Returns the number of tokens in a given text for a specific model.
    """
    encoding = tiktoken.encoding_for_model(model)
    tokens = encoding.encode(text)
    return len(tokens)
```

			price per token	value
Per review	input tokens	28063	0.0000025	0.0701575
	output tokens	2313	0.00001	0.02313
			total	0.09

<https://platform.openai.com/docs/pricing>

# Basic Prompt

The sky is

Complete the sentence:

The sky is

This is awesome! // Positive

This is bad! // Negative

Wow that movie was rad! // Positive

What a horrible show! //

# Core Components of an Effective Prompt

- **Instruction**

What you want the model to do — the **task or command**.

- **Context**

Additional information that helps guide the model toward a **better, more accurate response**.

- **Input Data**

The specific **question, text, or content** the model must analyze or respond to.

- **Output Indicator**

A description of the **desired format or style** of the response (e.g., "list," "JSON," "summary").

# Core Components of an Effective Prompt

When writing prompts, using **explicit variable names** (e.g., `<input>`, `<question>`, `<context>`) helps the model understand exactly what each part of the prompt represents.

Why this is a best practice:

- Makes prompts more structured and readable
- Reduces ambiguity for the model
- Helps maintain consistency when reusing or automating prompts
- Ideal for templates used in workflows or agent systems

Clear variable labels make prompts modular, reusable, and less error-prone.

# Zero-Shot Prompting

**Without providing any examples.** The model relies solely on the **instruction** and its **pre-trained knowledge** to generate the answer.

When to use it:

- For well-defined tasks the model already understands
- When you want fast, simple prompts
- When examples are unnecessary or could **bias the response**

Advantages:

- Easy to write
- Efficient
- Works well for many **general-purpose tasks**

Ask beforehand if the model knows what they're doing and understands what they need to do.

# Few-Shot Prompting

Provide the model with **a small number of examples** demonstrating how a task should be performed before giving the final query. These examples help the model infer the pattern, style, format, or reasoning approach you expect.

When to use it:

- When **Zero-Shot results are inconsistent**
- When the task requires a specific format or tone
- When you want the model to mimic a pattern or reasoning style

Advantages:

- More reliable outputs
- Better control over format and behavior
- Reduces ambiguity in complex tasks

# Few-Shot Prompting

*Prompt:*

```
This is awesome! // Negative  
This is bad! // Positive  
Wow that movie was rad! // Positive  
What a horrible show! //
```

*Output:*

```
Negative
```

# Chain-of-Thought Prompting

## Standard Prompting

Model Input

Q: Roger has 5 tennis balls. He buys 2 more cans of tennis balls. Each can has 3 tennis balls. How many tennis balls does he have now?

A: The answer is 11.

Q: The cafeteria had 23 apples. If they used 20 to make lunch and bought 6 more, how many apples do they have?

Model Output

A: The answer is 27.

## Chain-of-Thought Prompting

Model Input

Q: Roger has 5 tennis balls. He buys 2 more cans of tennis balls. Each can has 3 tennis balls. How many tennis balls does he have now?

A: Roger started with 5 balls. 2 cans of 3 tennis balls each is 6 tennis balls.  $5 + 6 = 11$ . The answer is 11.

Q: The cafeteria had 23 apples. If they used 20 to make lunch and bought 6 more, how many apples do they have?

Model Output

A: The cafeteria had 23 apples originally. They used 20 to make lunch. So they had  $23 - 20 = 3$ . They bought 6 more apples, so they have  $3 + 6 = 9$ . The answer is 9.

# Core Components of an Effective Prompt

Large tasks often require **multiple reasoning steps**, which increases the chance of errors.

By breaking the task into **smaller, clearer subtasks**, you:

- Reduce the **reasoning load** on the model
- Make each step simpler and more precise
- Improve the **accuracy** and **consistency** of the final result
- Allow the model to follow a **structured path** instead of guessing the entire solution at once

## **Key Idea:**

Smaller steps = Less ambiguity = Better outcomes.

# Agent-Based Approach

An **agent-based approach** uses multiple specialized “agents” — each with its own role, skills, or objectives — to solve a task collaboratively.


Instead of relying on a single prompt, the system coordinates **separate LLM-driven components** that interact, plan, and refine solutions.

Key characteristics:


- Specialized agents with distinct functions
- Communication between agents
- Iterative refinement and decision-making
- Modular and scalable system design

# Agent-Based Approach

## Single Agent



## Multi-Agent



[https://albanna-tutorials.com/llm\\_agents.html](https://albanna-tutorials.com/llm_agents.html)

# Prompt stability

## Prompt Stability Scoring for Text Annotation with Large Language Models

Christopher Barrie,<sup>1</sup> Elli Palaiologou,<sup>2</sup> Petter Törnberg<sup>3</sup>

<sup>1</sup>Department of Sociology, New York University

<sup>2</sup>Independent Researcher

<sup>3</sup>Institute for Logic, Language, and Computation, University of Amsterdam

February 18, 2025

<https://arxiv.org/pdf/2407.02039>

# Obrigado!

<https://www.promptingguide.ai>